

P4: Train a Smart Cab to Drive

Implement a basic driving agent

Our first policy was to choose a random move/action from (`None`, `'forward'`, `'left'`, `'right'`). Then we ran this agent within the simulation environment with `enforce_deadline` set to `False`.

In your report, mention what you see in the agent's behavior. Does it eventually make it to the target location?

The agent moves randomly. If I waited long enough, it would probably eventually make it to the target location, but I never saw it do so.

Identify and update state

Justify why you picked these set of states, and how they model the agent and its environment.

I first picked the following states:

- left
- right
- oncoming
- light

These four states seemed like a good start as they give enough information that the smart cab can go through intersections without breaking any traffic laws.

Implement Q-Learning

Using the Udacity Reinforcement Learning lectures as a reference, I implemented Q-learning. At this point, I also added one more item to the agent's state:

- next_waypoint

Adding this state allows the smart cab to make progress towards the goal. My agent seemed a bit more purposeful at this point. It did reach the goal several times (with `enforce_deadline` still set to `False`) as I watched it.

At this point, I decided to start recording the exact results. I now had the following settings:

- state: left, right, oncoming, light, next_waypoint
- discount (gamma): 0.9
- step size (alpha): 0.2

With these settings, the agent reached the primary destination 6 out of 10 times. Moreover, at least in the cases where the destination was reached, the cumulative reward was always positive.

Enhance the driving agent

Report what changes you made to your basic implementation of Q-Learning to achieve the final version of the agent. How well does it perform?

I considered a trial to have succeeded if it reached the destination and the cumulative reward was non-negative.

It took a lot of work on my implementation before I achieved good results. Without really thinking much about it, I first used a greedy policy for choosing the next action. I fiddled around with the value of gamma quite a lot, but could not do better than a success rate of about 40% (I must have gotten lucky with that first run success rate of 60%).

(Incidentally, at this point I needed to create some automation to summarize my results. With 100 trials, I could no longer count the success rate manually every time. My automation is in [analyze_data.py](#). It simply parses the output, which I would redirect to a file, and then creates a [spreadsheet](#).)

After spending many hours doing this, I finally went back and listened to the lectures again (always running experiments in the background). Watching the lectures again was not a bad idea. Using the ideas in the lecture, I implemented an epsilon-greedy strategy. Although I'm sure this probably improved things, it was not evident for 100 trials. I next tried an epsilon-decreasing strategy. Again, this did not seem to help (at least with 100 trials).

Finally, I tried adding "optimism in the face of uncertainty" initialization of the Q table. This made all the difference in the world. I was now achieving success rates of 80% and higher.

To achieve the final version of the agent, I experimented with different values of gamma. I recorded the results in this [spreadsheet](#).

With gamma set to 0.900, I achieved a success rate of 97%.

Does your agent get close to finding an optimal policy, i.e. reach the destination in the minimum possible time, and not incur any penalties?

I measured how close I was to finding an optimal policy by dividing the total rewards accrued during a trial by the maximum possible total rewards possible during that trial. You can see convergence towards an optimal policy in the plot below.

There is some fluctuation from trial to trial, but you can see that our policy converged to a value that was about 0.8 the value of the optimal policy.




Figure 1: learning growth